


The wobble is in the rubric, not the AI


We scored the same classroom observations over and over with the same AI, the same rubric and the same wording, then measured how often the verdict changed. Nothing varied between runs except the scorer itself, so any difference is the scorer’s.

The scorer proved consistent. Judgements were identical across 91% of re-scorings, and reliability held at 0.87 on a 0–1 scale — the range normally treated as dependable. The instability that remains is not spread evenly across the framework. It is concentrated in a handful of indicators whose wording admits two defensible readings of the same lesson.



Why those seven fail — two causes, two different fixes

Cause	Indicators	What is happening, and what fixes it
The standard bundles two requirements	C4, B2, C6, B9	A real lesson meets one and fails the other, and the rubric never says which decides. C4 asks for “deliberate strategies... diverse students included.” Runs answering yes cited the teacher calling on students by name; runs answering no cited that those were volunteers while many stayed silent. Both readings are correct. Rewrite: split the standard, or state which clause governs.
The evidence is not in the audio	D4, D5, D6	Sustained focus during silent work, willingness to attempt, whether pupils handle materials — none of this is audible. The scorer is inferring from silence. No rewrite helps. Needs video, or reword around what can be heard.



What to do

1. Rewrite the four compound standards — C4, B2, C6, B9. Highest-leverage change available, and it costs nothing but editing time.
2. Decide about D4, D5 and D6. Either accept they need video, or reword them around what an audio transcript can actually evidence.
3. Report the overall lesson score from a single AI pass. Do not report Section B or Section D from one pass — their High/Medium/Low band moves between re-scorings on about half of lessons.

Two limits. This measures whether the scorer agrees with itself, not whether it is right — an indicator returning the same wrong answer ten times scores perfectly here. And the failures cannot be spotted by reading the rubric: across all 37 indicators no textual feature correlated with the disagreement actually observed (all $|p| < 0.16$). Which indicators fail has to be measured, not inspected.